

## README Programming Assignment 2 Winter 2015

Author: Dinesh Jayasankar

Problem #1 Demonstrate with F7 key.

### **/usr/src/servers/is/proto.h**

In proto.h, I declared \_PROTOTYPE( void magic\_dmp, (void) );

### **/usr/src/servers/is/dmp.c**

I assigned magic\_dmp to function key F7 in struct hook\_entry hooks[NHOOKS]

### **/usr/src/servers/is/dmp\_kernel.c**

I defined PUBLIC in messages[NR\_TASKS + NR\_PROCS][NR\_TASKS + NR\_PROCS];

I define PUBLIC void magic\_dmp(). Inside magic\_dmp(), I retrieved messages matrix with sys\_getmagic(messages). The outer for-loop prints the first all the processes 10 at a time. In the inner for-loop, for each process printed, I print the first 10 destination processes that receive a message. If the source or destination process is shutdown, which is checked with isempty(), I skip that iteration of the loop.

### **/usr/src/kernel/proc.h**

I declared the matrix EXTERN messages[NR\_TASKS + NR\_PROCS][NR\_TASKS + NR\_PROCS].

### **/usr/src/include/minix/syslib.h**

I declare sys\_getmagic(dst) as a shorthand for sys\_getinfo(GET\_MAGIC, dst, 0, 0, 0)

It gets a copy of system info or kernel data.

### **/lib/syslib/sys\_getinfo.c**

This contains PUBLIC do\_getinfo(...)

### **/usr/src/kernel/system/do\_getinfo.c**

This contains PUBLIC int do\_getinfo(m\_ptr).

I have a case statement "case GET\_MAGIC". The length is initialized to sizeof(int) \*

(NR\_PROCS+NR\_TASKS) \* (NR\_PROCS+NR\_TASKS). I use vir2phys() to get the physical bytes of the integer message matrix.

### **/usr/src/kernel/proc.c**

I increment the appropriate entry in the two-dimensional message matrix inside mini\_send(). The row index is caller\_ptr->p\_nr + NR\_TASKS, and column index is dst\_ptr->p\_nr + NR\_TASKS

### **include/minix/com.h**

NR\_TASKS is defined here and IDLE is -4. Since Kernel tasks have negative process numbers, we require the NR\_TASKS offset when indexing the message matrix.

Problem #3 Demonstrate with F11 key

**/include/minix/callnr.h (page 676).**

#define NRCALLS 91 is the number of system calls allowed. The user program's typical system calls range from 1 to 63.

**/usr/src/kernel/proc.h**

I declared the fields "int calls[NRCALLS]" to record each instance of a system call and "int num\_of\_calls" to record system call frequency.

**/usr/src/other/syscall.c**

I modified \_sendrec(...) to accept three parameters: who, msgptr, syscallnr

The function \_syscall handles all system calls from user programs.

\_syscall takes as parameters who, syscallnr, and msgptr.

Example:

- ❖ if (fork() != 0)
- ❖ /lib/syscall/fork.s
  - ... . \_fork: jmp \_fork defined in src/lib/posix/fork.c
- ❖ /src/lib/posix/fork.c
  - ... return(\_syscall(PM, FORK, &m));

\_syscall calls \_sendrec(who, msgptr).

NOTE: \_sendrec calls the assembly code in lib/i386/rts/\_ipc.s or lib/i86/rts/\_sendrec.s

**/usr/src/include/minix/ipc.h and /include/minix/ipc.h**

I changed \_sendrec's prototype to \_PROTOTYPE( int sendrec, (int src\_dst, message \*m\_ptr, ...));

Note the default arguments. This is possible due to #include <stdarg.h> (variadic functions).

**/lib/i386/rts/\_ipc.s and /lib/i86/rts/\_sendrec.s**

NOTE: Documentation says that edx is unused.

**\_ipc.s:**

I added line "mov edx, CALL\_NR(ebp)" to store the extra syscallnr parameter.. CALL\_NR is 16, because the parameter syscallnr comes after msgptr.

**\_sendrec.s**

I added the line "mov dx, 8(bp)" to store the extra syscallnr parameter. Note that this is not using extended register notation, so the offset different will be by 2, instead of 4.

**/kernel/mpx386.s (\_s\_call)**

NOTE: The register edi is unused, but the parameter is in edx on the stack.

Change all references of dx to di.

mov di, ss

mov ds, di

mov es, di

Place edx in the right place using “push edx”. Now, sys\_call can access this as a normal parameter.

```
xor     ebp, ebp      ! for stacktrace
                        ! end of inline save
                        ! now set up parameters for sys_call
push     edx
push     ebx          ! pointer to user message
push     eax          ! src/dest
push     ecx          ! SEND/RECEIVE/BOTH
call     _sys_call    ! sys_call(function, src_dest,
                        ! caller is now explicitly in
```

#### **/kernel/proto.h**

Modify the parameters of syscall to include call\_index which will hold the user-space value of syscallnr.

#### **/kernel/proc.c**

If the system call number is between 1 and 63, the function is a SENDREC, and proc\_ptr is a user process then use call\_index (syscallnr) to index the field “calls” and increment it.

#### **/kernel/system/do\_fork.c**

On a new process fork reset the calls and num\_of\_calls field!

Problem #2 Demonstrate with F1 key.

#### **/kernel/proc.c (sched)**

If it is a user process and its system and user time are 0, then set “time\_left” to 1 and give the new process a new quantum. A fork inherits the parent's existing p\_ticks\_left.

This sets \*front = time\_left -> 1.

#### **/kernel/proc.c enqueue**

Since front is 1, the new process will be placed at the front of the queue.